

Lab-1 Writeup

Enhancement 1: The current code does not actually evaluate the model on the test set, but it only evaluates it on the val set. When you write papers, you would ideally split the dataset into train, val and test. Train and val are both used in training, and the model trained on the training data, and evaluated on the val data. So why do we need test split? We report our results on the test split in papers. Also, we do cross-validation on the train/val split (covered in later labs).

Report the results of the model on the test split. (Hint: It would be exactly like the evaluation on the val dataset, except it would be done on the test dataset.)

- We need a test split to avoid overfitting on the Val-dataset.
- Below are the lines of code I updated.

- Import the relevant metrics.

```
from sklearn.metrics import (
    accuracy_score,
    # ENHANCEMENT 1
    precision_score,
    recall_score,
    f1_score
)
```

- Run up a tally and get the avg metric per epoch.

```
def evaluate(model, data_loader, criterion):
    epoch_loss = 0
    epoch_acc = 0
    # ENHANCEMENT 1
    epoch_precision = 0
    epoch_recall = 0
    epoch_f1 = 0

    # ENHANCEMENT 1
    precision = precision_score(labels.cpu(), preds.cpu(), average='weighted', zero_division=0)
    recall = recall_score(labels.cpu(), preds.cpu(), average='weighted', zero_division=0)
    f1 = f1_score(labels.cpu(), preds.cpu(), average='weighted', zero_division=0)
    epoch_precision += precision
    epoch_recall += recall
    epoch_f1 += f1

    num_batches = len(data_loader)
    return (
        epoch_loss / num_batches,
        epoch_acc / num_batches,

        # ENHANCEMENT 1
        epoch_precision / num_batches,
        epoch_recall / num_batches,
        epoch_f1 / num_batches
    )
```

- Update unpacking and print statement in training (note that some is cut off).

```
# Let's train our model
for epoch in range(100):
    train_loss, train_acc = train(winemodel, train_dataloader, optimizer, criterion)
    # ENHANCEMENT 1
    valid_loss, valid_acc, valid_precision, valid_recall, valid_f1 = evaluate(winemodel, val_dataloader, criterion)

    print(f'| Epoch: {epoch+1:02} | Train Loss: {train_loss:.3f} | Train Acc: {train_acc*100:.2f}% | Val. Loss: {valid_loss:.3f}
```

- Run evaluate on test data and update print statement (note that some is cut off).

```
test_loss, test_acc, test_prec, test_rec, test_f1 = evaluate(winemodel, test_dataloader, criterion)
print(f'| Test. Loss: {test_loss:.3f} | Test. Acc: {test_acc*100:.2f}% | Test. Precision: {test_prec:.3f} | Test. Recall:
```

- Train-Dev Results
 - | Epoch: 100 | Train Loss: 1.302 | Train Acc: 74.92% | Val. Loss: 1.442 | Val. Acc: 58.13% | Val. Precision: 0.661 | Val. Recall: 0.581 | Val. F1-score: 0.590 |
- Test-Results
 - | Test. Loss: 1.422 | Test. Acc: 62.29% | Test. Precision: 0.702 | Test. Recall: 0.623 | Test. F1-score: 0.626 |

Enhancement 2: Increase the number of epochs (and maybe the learning rate).

RESULTS

- Code I updated:


```
# ENHANCEMENT 2
for epoch in range(500):
    # Define and the Loss function and optimizer
    criterion = nn.CrossEntropyLoss().to(device)
    # ENHANCEMENT 2
    optimizer = AdamW(winemodel.parameters(), lr = 1e-2)
```
- Train-Dev Results
 - | Epoch: 500 | Train Loss: 1.629 | Train Acc: 41.41% | Val. Loss: 1.575 | Val. Acc: 46.88% | Val. Precision: 0.289 | Val. Recall: 0.468 | Val. F1-score: 0.347 |
- Test Results
 - | Test. Loss: 1.575 | Test. Acc: 46.88% | Test. Precision: 0.277 | Test. Recall: 0.469 | Test. F1-score: 0.337 |

Does the accuracy on the test set increase?

- NO! It was 15.81% worse. Down from 62.69% to 46.88%.

Is there a significant difference between the test accuracy and the train accuracy? If yes, why?

- The test-accuracy was better than the final epoch of the train-accuracy by about five percent. I think that this is because the learning rate is too large. The model performed best on the original learning rate: 1e-3. I tried larger (1e-2, 1e-1) and smaller (1e-4) learning rates but the original, regardless of the number of epochs (up to 500), has performed the best. Sometimes larger learning rates perform much worse than smaller learning rates because they can overshoot the true minimum repeatedly.

Enhancement 3: Increase the depth of your model (add more layers). Report the parts of the model definition you had to update. Report results.

- CODE I UPDATED

```
def __init__(self):
    super(WineModel, self).__init__()

    self.linear1 = torch.nn.Linear(11, 200)
    self.activation = torch.nn.ReLU()

    # ENHANCEMENT 3
    self.linear2 = torch.nn.Linear(200, 200)
    self.linear3 = torch.nn.Linear(200, 6)

    self.softmax = torch.nn.Softmax(dim =1)
```

```
def forward(self, x):
    x = self.linear1(x)
    x = self.activation(x)
    x = self.linear2(x)

    # ENHANCEMENT 3
    x = self.activation(x)
    x = self.linear3(x)

    x = self.softmax(x)
    return x
```

- Train-Dev Results
 - | Epoch: 500 | Train Loss: 1.211 | Train Acc: 83.20% | Val. Loss: 1.447 | Val. Acc: 60.00% | Val. Precision: 0.573 | Val. Recall: 0.600 | Val. F1-score: 0.561 |
- Test Results
 - | Test. Loss: 1.375 | Test. Acc: 66.88% | Test. Precision: 0.648 | Test. Recall: 0.669 | Test. F1-score: 0.633 |

Enhancement 4: Increase the width of your model's layers. Report the parts of the model definition you had to update. Report results.

```
class WineModel(torch.nn.Module):

    def __init__(self):
        super(WineModel, self).__init__()
        self.linear1 = torch.nn.Linear(11, 200) # ENHANCEMENT 4 (Larger right number)
        self.activation = torch.nn.ReLU()
        self.linear2 = torch.nn.Linear(200, 400) # ENHANCEMENT 4 (Larger numbers)
        self.linear3 = torch.nn.Linear(400, 6) # ENHANCEMENT 4 (Larger Left number)
        self.softmax = torch.nn.Softmax(dim =1)
```

- Train-Dev Results
 - | Epoch: 500 | Train Loss: 1.234 | Train Acc: 81.02% | Val. Loss: 1.367 | Val. Acc: 67.50% | Val. Precision: .6600 | Val. Recall: .6750 | Val. F1-score: .6367 |
- Test Results
 - | Test. Loss: 1.461 | Test. Acc: 56.67% | Test. Precision: 0.586 | Test. Recall: 0.567 | Test. F1-score: 0.539 |

Enhancement 5: Choose a new dataset from the list below. Search the Internet and download your chosen dataset (many of them could be available on kaggle). Adapt your model to your dataset. Train your model and record your results.

- | Epoch: 50 | Train Loss: 0.318 | Train Acc: 99.56% | Val. Loss: 0.325 | Val. Acc: 98.33% | Val. Precision: 1.000 | Val. Recall: 0.983 | Val. F1-score: 0.990 |
- | Test. Loss: 0.349 | Test. Acc: 96.43% | Test. Precision: 0.979 | Test. Recall: 0.964 | Test. F1-score: 0.964 |

Enhancement 5: Choose a new dataset from the list below. Search the Internet and download your chosen dataset (many of them could be available on kaggle). Adapt your model to your dataset. Train your model and record your results.

- cancer_dataset - Breast cancer dataset.
- crab_dataset - Crab gender dataset.
- glass_dataset - Glass chemical dataset.
- iris_dataset - Iris flower dataset.
- ovarian_dataset - Ovarian cancer dataset.
- thyroid_dataset - Thyroid function dataset.

Imports

```
In [47]: import pandas as pd
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.optim import AdamW
from torch.utils.data import Dataset, DataLoader
from tqdm.notebook import tqdm
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

Load Dataset

```
In [48]: from sklearn.datasets import load_breast_cancer
cancer_dataset = load_breast_cancer()
```

```
In [49]: type(cancer_dataset)
```

```
Out[49]: sklearn.utils.Bunch
```

```
In [50]: df = pd.DataFrame(data=cancer_dataset.data, columns=cancer_dataset.feature_names)
df['target'] = cancer_dataset.target
```

```
In [51]: # make sure it looks good
df.head()
# df.tail()
```

Out[51]:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	target
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419	0
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812	0
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069	0
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597	0
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809	0

5 rows × 11 columns

```
In [52]: # how many features?  
# WOW!  
len(df.columns) - 1
```

Out[52]: 30

```
In [53]: # 0 means malignant  
# 1 means benign  
# @$$ switch these  
df['target'].tail(5)
```

```
Out[53]: 564    0  
565    0  
566    0  
567    0  
568    1  
Name: target, dtype: int32
```

```
In [54]: # normalize data so that each feature has a mean of 0 and a std of 1  
labels = df['target']  
# @$$ there is a better way  
# if vals too big. -> maybe each col normalized seperatly  
df = (df - df.mean()) / df.std()  
df['target'] = labels
```

```
In [55]: df.head()
```

Out[55]:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	me symmet
0	1.096100	-2.071512	1.268817	0.983510	1.567087	3.280628	2.650542	2.530249	2.2155
1	1.828212	-0.353322	1.684473	1.907030	-0.826235	-0.486643	-0.023825	0.547662	0.0013
2	1.578499	0.455786	1.565126	1.557513	0.941382	1.052000	1.362280	2.035440	0.9388
3	-0.768233	0.253509	-0.592166	-0.763792	3.280667	3.399917	1.914213	1.450431	2.8648
4	1.748758	-1.150804	1.775011	1.824624	0.280125	0.538866	1.369806	1.427237	-0.0095

5 rows × 31 columns

In [56]: `# sumamry statistics of the data
df.describe()`

Out[56]:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	
count	5.690000e+02	5.690000e+02	5.690000e+02	5.690000e+02	5.690000e+02	5.690000e+02	5.6
mean	-3.142575e-15	-6.558316e-15	-7.012551e-16	-8.339355e-16	6.083788e-15	-1.081346e-15	-3.
std	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.0
min	-2.027864e+00	-2.227289e+00	-1.982759e+00	-1.453164e+00	-3.109349e+00	-1.608721e+00	-1.
25%	-6.887793e-01	-7.253249e-01	-6.913472e-01	-6.666089e-01	-7.103378e-01	-7.464292e-01	-7.
50%	-2.148925e-01	-1.045442e-01	-2.357726e-01	-2.949274e-01	-3.486040e-02	-2.217454e-01	-3.
75%	4.689800e-01	5.836621e-01	4.992377e-01	3.631877e-01	6.356397e-01	4.934227e-01	5.
max	3.967796e+00	4.647799e+00	3.972634e+00	5.245913e+00	4.766717e+00	4.564409e+00	4.2

8 rows × 31 columns

Load this dataset for training a neural network

```
In [57]: # The dataset class  
# inherit the Dataset class and create a new class CancerDataset  
class CancerDataset(Dataset):  
    def __init__(self, df):  
        self.df = df  
        self.features = []  
        self.labels = []  
  
        # Iterate through rows of the DataFrame  
        for _, row in df.iterrows():  
            # Append features for each row except the 'target' column  
            self.features.append(row.drop('target').tolist())  
            # Append label for each row  
            self.labels.append(row['target'])
```

```

# set length to return number of rows
def __len__(self):
    return len(self.df)

# get a sample in the form of a dictionary (features and labels) given its index
def __getitem__(self, idx):
    if torch.is_tensor(idx):
        idx = idx.tolist()

    features = self.features[idx]
    features = torch.FloatTensor(features)

    labels = torch.tensor(int(self.labels[idx]), dtype = torch.long)

    return {'labels': labels, 'features': features}

# instantiate a CancerDataset object based off of data_df
cancer_dataset = CancerDataset(df)
# perform a 80, 10, 10 split using the cancer_dataset object
train_dataset, val_dataset, test_dataset = torch.utils.data.random_split(cancer_dataset, [80, 10, 10])

# The dataloader
train_dataloader = DataLoader(
    # dataset input
    train_dataset,
    # batches of 4 during training
    batch_size = 4,
    # don't learn from the order of the dataset!
    shuffle = True,
    # multy processing. (0 = 1 process to load data)
    # how many processes can you do?
    num_workers = 0
)

#want results without any randomness, therefore shuffle is False for these
val_dataloader = DataLoader(val_dataset, batch_size = 4, shuffle = False, num_workers = 0)
test_dataloader = DataLoader(test_dataset, batch_size = 4, shuffle = False, num_workers = 0)

```

```

In [58]: # peak into the dataset
for i in cancer_dataset:
    print(i)
    break

```

```

{'labels': tensor(0), 'features': tensor([ 1.0961, -2.0715,  1.2688,  0.9835,  1.567
1,  3.2806,  2.6505,  2.5302,
        2.2156,  2.2538,  2.4875, -0.5648,  2.8305,  2.4854, -0.2138,  1.3157,
        0.7234,  0.6602,  1.1477,  0.9063,  1.8850, -1.3581,  2.3016,  1.9995,
        1.3065,  2.6144,  2.1077,  2.2941,  2.7482,  1.9353])}

```

Neural Network

```

In [59]: # change the device to gpu if available
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

```

```

In [60]: #defines new class named CancerModel, inherits from torch.nn.Module
class CancerModel(torch.nn.Module):

```

```

def __init__(self):
    # initialize in the same way the parent class does first
    super(CancerModel, self).__init__()

    # create linear layer. 11 inputs 200 outputs?
    self.linear1 = torch.nn.Linear(30, 210)
    # creates activation function (What does this do and how does it work @$$)
    self.activation = torch.nn.ReLU()

    # a second linear layer. Takes in 200 and outputs 6 (one corresponding to each
    self.linear2 = torch.nn.Linear(210, 2)

    # sigmoid activation for binary
    self.softmax = torch.nn.Softmax(dim=1)
    # defines forward pass of nn
    # sends input through the tunnel: lin1 -> act -> line2 -> softmax
    def forward(self, x):
        x = self.linear1(x)
        x = self.activation(x)
        x = self.linear2(x)
        x = self.activation(x)
        x = self.softmax(x)
        return x

    # earlier we set device to cpu or gpu
    # create this CancerModel object, then move to cpu or gpu
    cancermodel = CancerModel().to(device)

```

```

In [61]: # Define and the loss function and optimizer
          # we are trying binary cross entropy loss
          # what is BCE using as inputs
          #
          criterion = nn.CrossEntropyLoss().to(device)

          # an optimizer @$$ what does it do?
          # specify the learning rate
          optimizer = AdamW(cancermodel.parameters(), lr = 1e-3)

```

```

In [62]: # Lets define the training steps
def accuracy(preds, labels):
    preds = torch.argmax(preds, dim=1).flatten()
    return torch.sum(preds == labels) / len(labels)

def train(model, data_loader, optimizer, criterion):
    epoch_loss = 0
    epoch_acc = 0

    # activate training mode
    model.train()

    # iterate over batches
    for d in tqdm(data_loader):
        inputs = d['features'].to(device)
        #labels = d['labels'].unsqueeze(1).to(device) # Adjusting the shape of the ta
        labels = d['labels'].to(device) # Adjusting the shape of the target tensor

        #labels = labels.float() # Cast labels to float32

        # obtaining predictions

```



```

        outputs = cancermodel(inputs)
        preds = (outputs > 0.5).float() # Convert logits to binary predictions
        # calculate loss and accuracy by comparing predictions to labels
        loss = criterion(outputs, labels)
        acc = accuracy(preds, labels) # Use binary accuracy function

    # Backpropagation
    loss.backward()
    optimizer.step()
    optimizer.zero_grad()

    epoch_loss += loss.item()
    epoch_acc += acc.item()

    return epoch_loss / len(data_loader), epoch_acc / len(data_loader)

# Lets define the testing steps
def evaluate(model, data_loader, criterion):
    epoch_loss = 0
    epoch_acc = 0
    epoch_precision = 0
    epoch_recall = 0
    epoch_f1 = 0

    # eval mode activate!
    model.eval()

    # iterate over batches (later we will average the huge sum)
    with torch.no_grad():
        for d in data_loader:
            # extract labels and features from batch and send to cpu/gpu
            inputs = d['features'].to(device)
            labels = d['labels'].to(device) # Adjusting the shape of the target tensor
            # labels = labels.float() # Convert labels to float32

            # use model to get predictions
            outputs = cancermodel(inputs)

            # get loss and accuracy
            loss = criterion(outputs, labels)
            acc = accuracy(outputs, labels)

            # run a big sum!
            epoch_loss += loss.item()
            epoch_acc += acc.item()
            precision = precision_score(labels.cpu(), outputs.argmax(dim=1).cpu() > 0.5, average='weighted')
            recall = recall_score(labels.cpu(), outputs.argmax(dim=1).cpu() > 0.5, average='weighted')
            f1 = f1_score(labels.cpu(), outputs.argmax(dim=1).cpu() > 0.5, average='weighted')
            epoch_precision += precision
            epoch_recall += recall
            epoch_f1 += f1

    # average big sums by # of epochs
    num_batches = len(data_loader)
    return (epoch_loss / num_batches,
            epoch_acc / num_batches,
            epoch_precision / num_batches,
            epoch_recall / num_batches,
            epoch_f1 / num_batches)

```

```
# Let's train our model
for epoch in range(50):
    train_loss, train_acc = train(cancermodule, train_dataloader, optimizer, criterion)
    valid_loss, valid_acc, valid_precision, valid_recall, valid_f1 = evaluate(cancermodule, valid_dataloader)

    print(f'| Epoch: {epoch+1:02} | Train Loss: {train_loss:.3f} | Train Acc: {train_acc:.3f} | Valid Loss: {valid_loss:.3f} | Valid Acc: {valid_acc:.3f} | Valid Precision: {valid_precision:.3f} | Valid Recall: {valid_recall:.3f} | Valid F1: {valid_f1:.3f}')
```

[illegible]

[illegible]

```

| Epoch: 38 | Train Loss: 0.318 | Train Acc: 99.56% | Val. Loss: 0.347 | Val. Acc: 9
6.67% | Val. Precision: 0.978 | Val. Recall: 0.967 | Val. F1-score: 0.964 |
0%| | 0/114 [00:00<?, ?it/s]
| Epoch: 39 | Train Loss: 0.318 | Train Acc: 99.56% | Val. Loss: 0.347 | Val. Acc: 9
6.67% | Val. Precision: 0.978 | Val. Recall: 0.967 | Val. F1-score: 0.964 |
0%| | 0/114 [00:00<?, ?it/s]
| Epoch: 40 | Train Loss: 0.318 | Train Acc: 99.56% | Val. Loss: 0.347 | Val. Acc: 9
6.67% | Val. Precision: 0.978 | Val. Recall: 0.967 | Val. F1-score: 0.964 |
0%| | 0/114 [00:00<?, ?it/s]
| Epoch: 41 | Train Loss: 0.318 | Train Acc: 99.56% | Val. Loss: 0.346 | Val. Acc: 9
6.67% | Val. Precision: 0.978 | Val. Recall: 0.967 | Val. F1-score: 0.964 |
0%| | 0/114 [00:00<?, ?it/s]
| Epoch: 42 | Train Loss: 0.318 | Train Acc: 99.56% | Val. Loss: 0.347 | Val. Acc: 9
6.67% | Val. Precision: 0.978 | Val. Recall: 0.967 | Val. F1-score: 0.964 |
0%| | 0/114 [00:00<?, ?it/s]
| Epoch: 43 | Train Loss: 0.318 | Train Acc: 99.56% | Val. Loss: 0.347 | Val. Acc: 9
6.67% | Val. Precision: 0.978 | Val. Recall: 0.967 | Val. F1-score: 0.964 |
0%| | 0/114 [00:00<?, ?it/s]
| Epoch: 44 | Train Loss: 0.318 | Train Acc: 99.56% | Val. Loss: 0.347 | Val. Acc: 9
6.67% | Val. Precision: 0.978 | Val. Recall: 0.967 | Val. F1-score: 0.964 |
0%| | 0/114 [00:00<?, ?it/s]
| Epoch: 45 | Train Loss: 0.318 | Train Acc: 99.56% | Val. Loss: 0.347 | Val. Acc: 9
6.67% | Val. Precision: 0.978 | Val. Recall: 0.967 | Val. F1-score: 0.964 |
0%| | 0/114 [00:00<?, ?it/s]
| Epoch: 46 | Train Loss: 0.318 | Train Acc: 99.56% | Val. Loss: 0.347 | Val. Acc: 9
6.67% | Val. Precision: 0.978 | Val. Recall: 0.967 | Val. F1-score: 0.964 |
0%| | 0/114 [00:00<?, ?it/s]
| Epoch: 47 | Train Loss: 0.318 | Train Acc: 99.56% | Val. Loss: 0.347 | Val. Acc: 9
6.67% | Val. Precision: 0.978 | Val. Recall: 0.967 | Val. F1-score: 0.964 |
0%| | 0/114 [00:00<?, ?it/s]
| Epoch: 48 | Train Loss: 0.318 | Train Acc: 99.56% | Val. Loss: 0.348 | Val. Acc: 9
6.67% | Val. Precision: 0.978 | Val. Recall: 0.967 | Val. F1-score: 0.964 |
0%| | 0/114 [00:00<?, ?it/s]
| Epoch: 49 | Train Loss: 0.318 | Train Acc: 99.56% | Val. Loss: 0.347 | Val. Acc: 9
6.67% | Val. Precision: 0.978 | Val. Recall: 0.967 | Val. F1-score: 0.964 |
0%| | 0/114 [00:00<?, ?it/s]
| Epoch: 50 | Train Loss: 0.318 | Train Acc: 99.56% | Val. Loss: 0.347 | Val. Acc: 9
6.67% | Val. Precision: 0.978 | Val. Recall: 0.967 | Val. F1-score: 0.964 |

```

```

In [64]: test_loss, test_acc, test_prec, test_rec, test_f1 = evaluate(cancermodel, test_data_loader)

print(f'| Test. Loss: {test_loss:.3f} | Test. Acc: {test_acc*100:.2f}% | Test. Precision: {test_prec:.3f} | Test. Recall: {test_rec:.3f} | Test. F1-score: {test_f1:.3f} |')

```